



Apache Flink

基于Flink on Kubernetes 的大数据平台

张凯 · 阿里云 / 高级技术专家
任春德 · 阿里云 / 高级技术专家

Apache Flink Meetup 北京 - 2019年06月29日



CONTENT

目录 >>

01 / K8s 简介

02 / Flink on K8s

03 / 大数据平台

04 / 展望



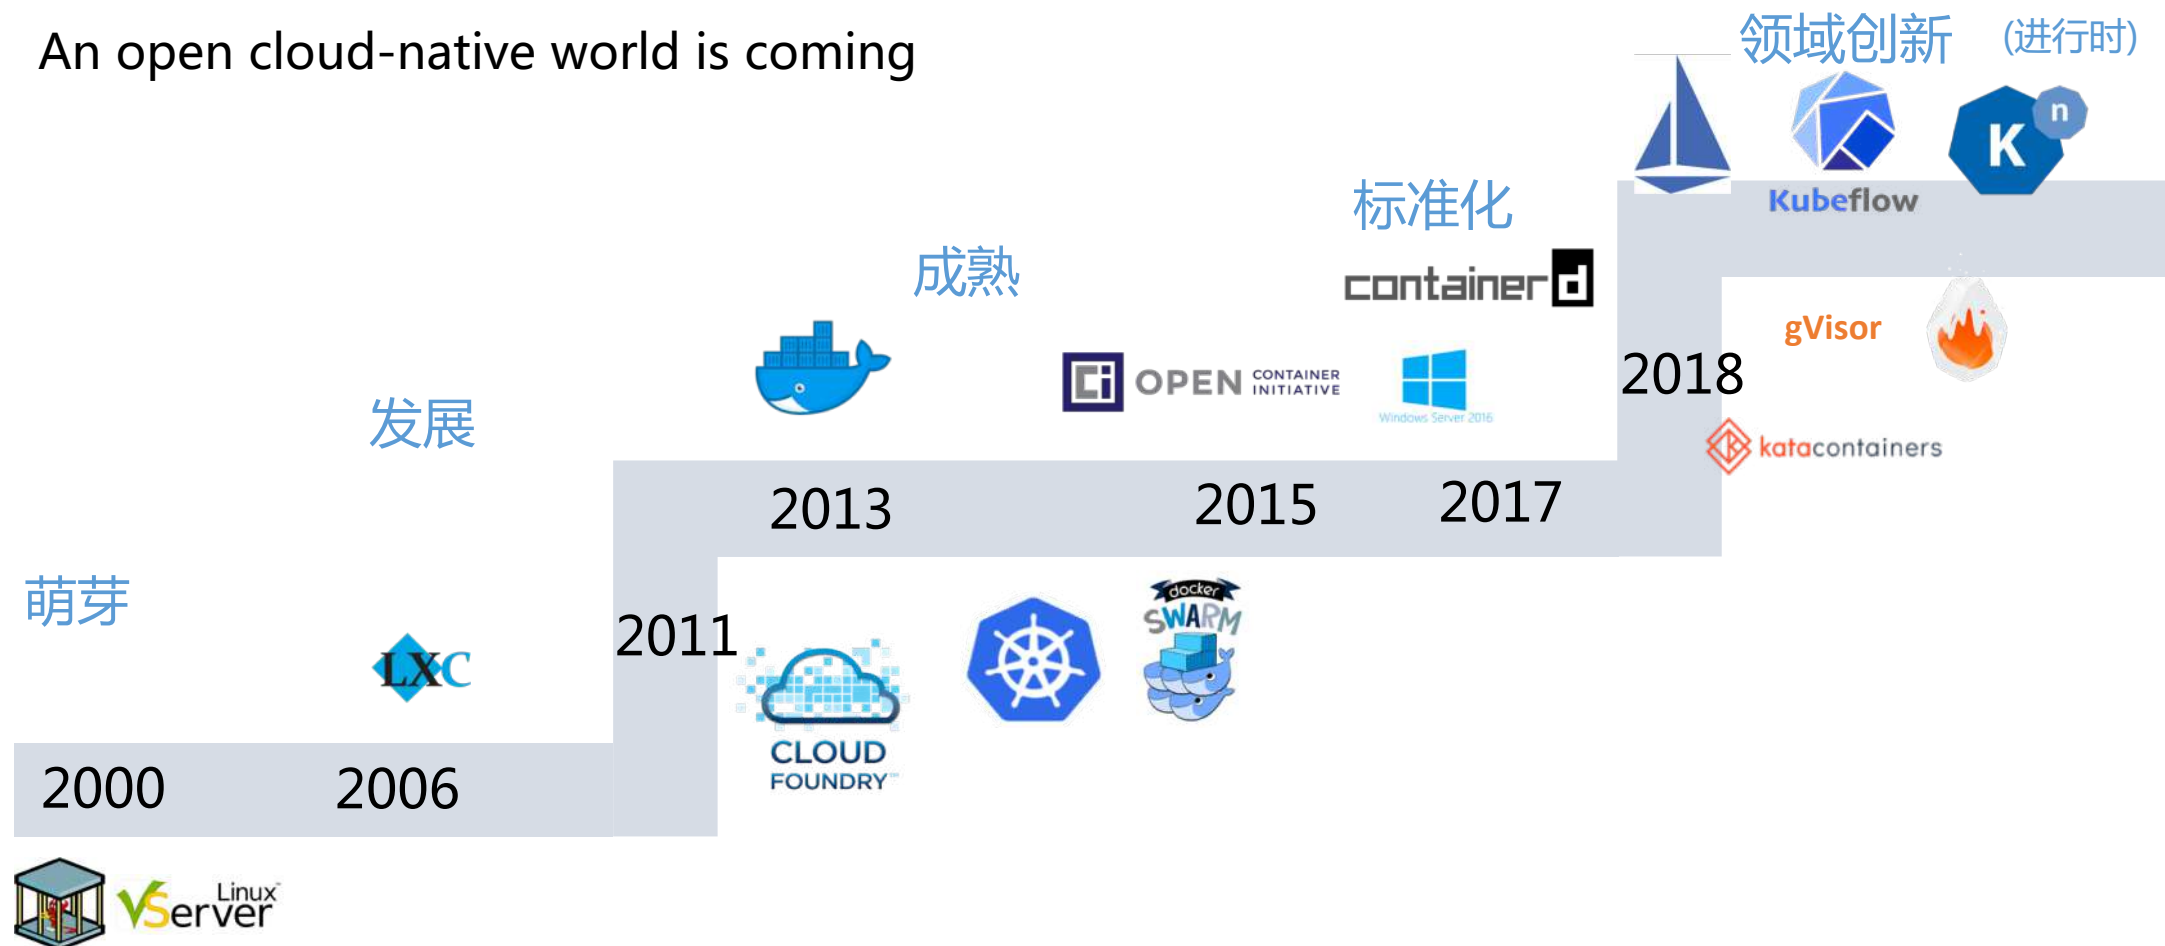
01

K8S简介



Evolution of Container and Cloud Native technology

An open cloud-native world is coming





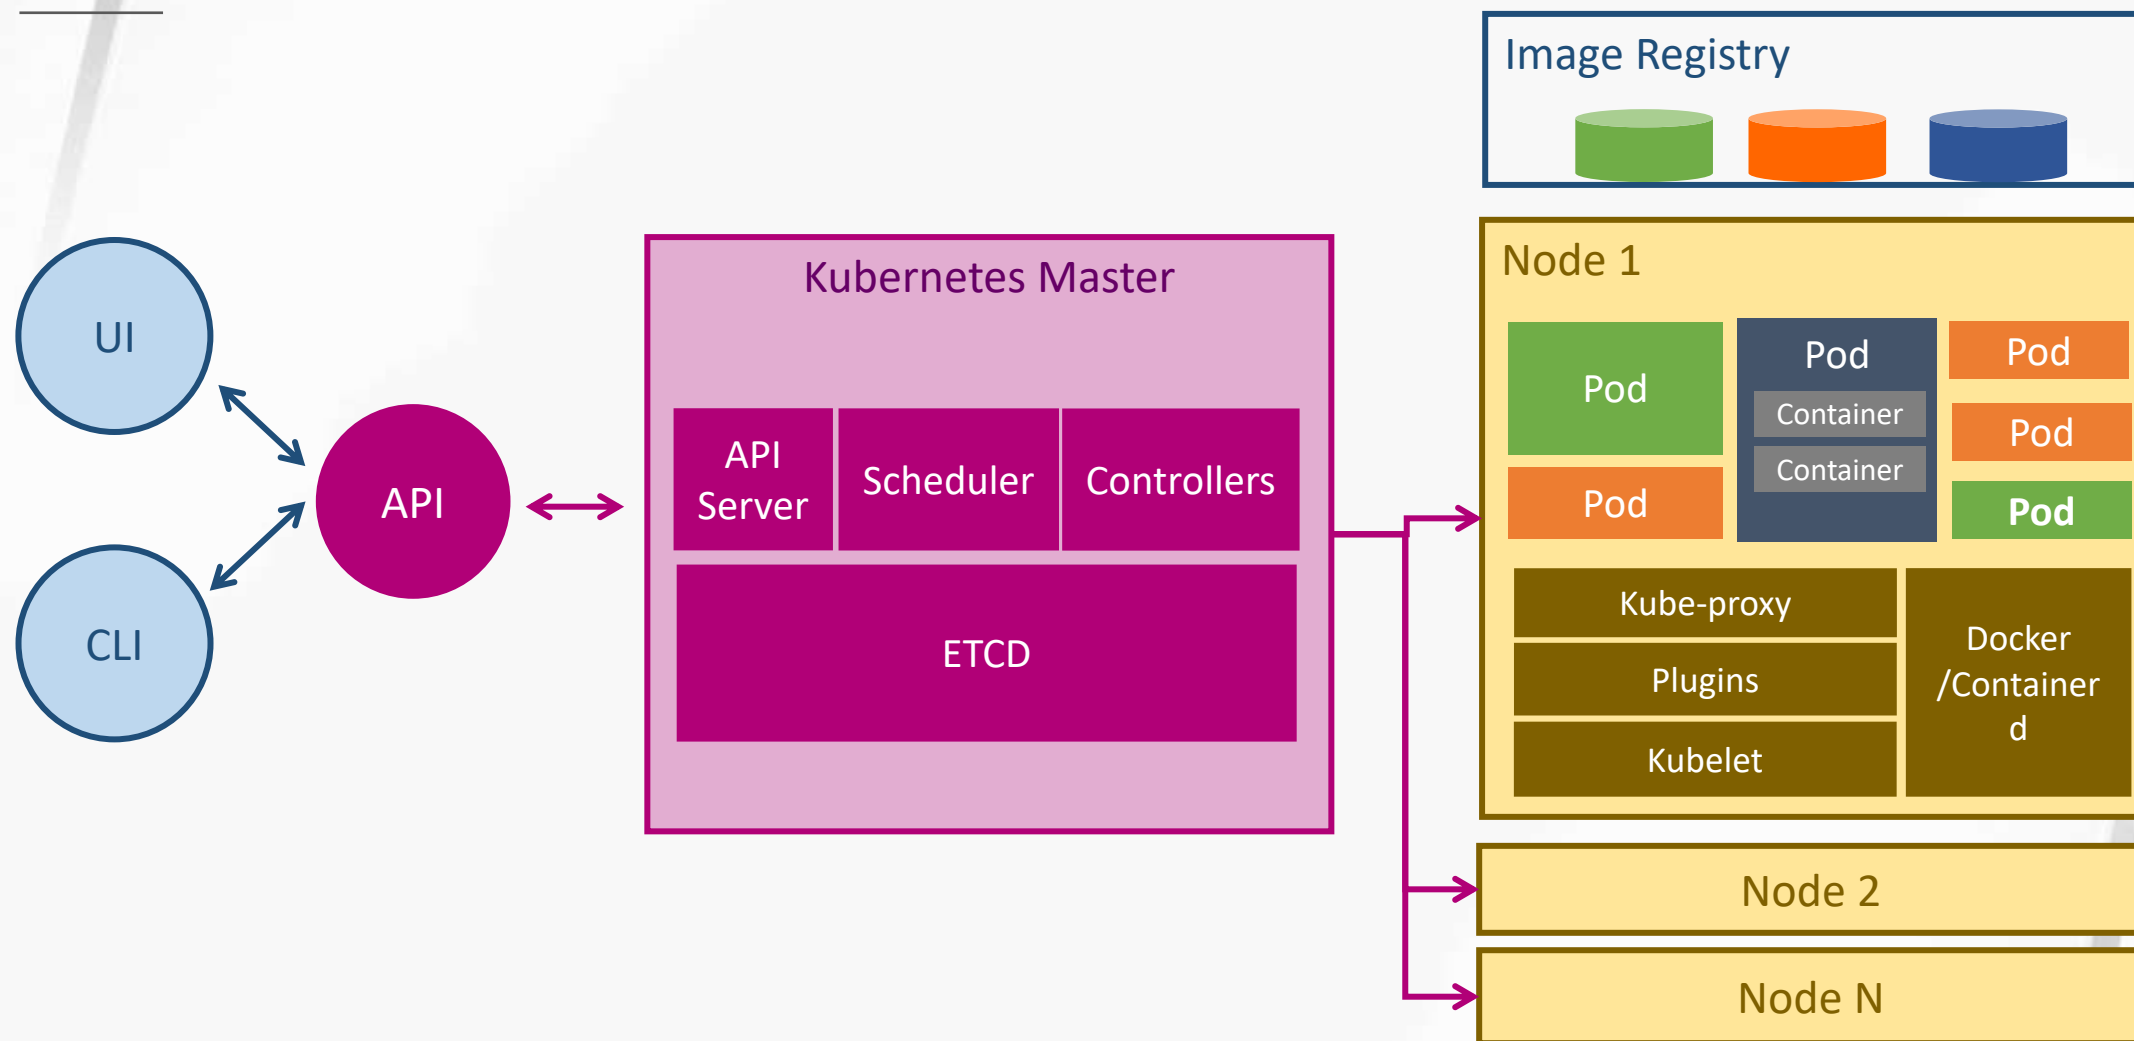
What is Kubernetes?

- Kubernetes is an open source container orchestration engine for automating deployment, scaling, and management of containerized applications.
- Resource management
- Container orchestration
- Application management
- Operation automation
- Cloud native application platform



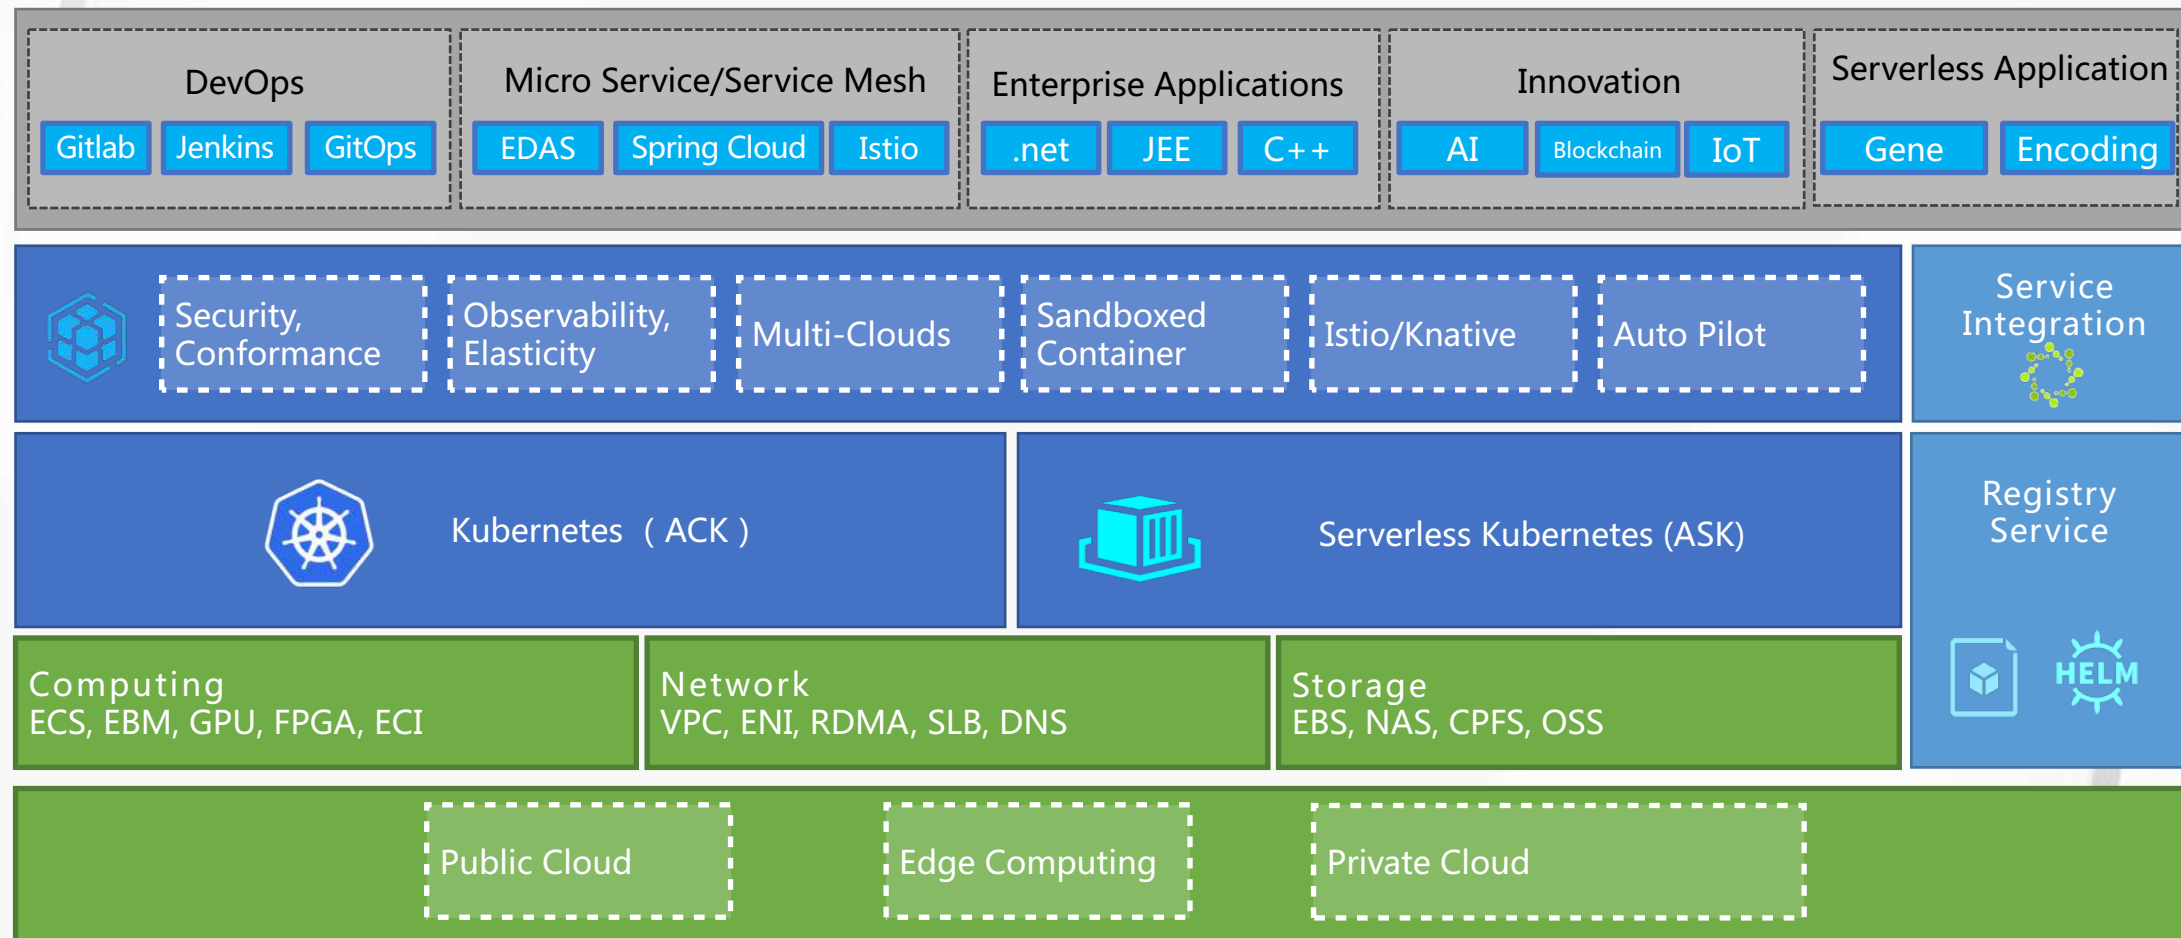


Kubernetes Architecture





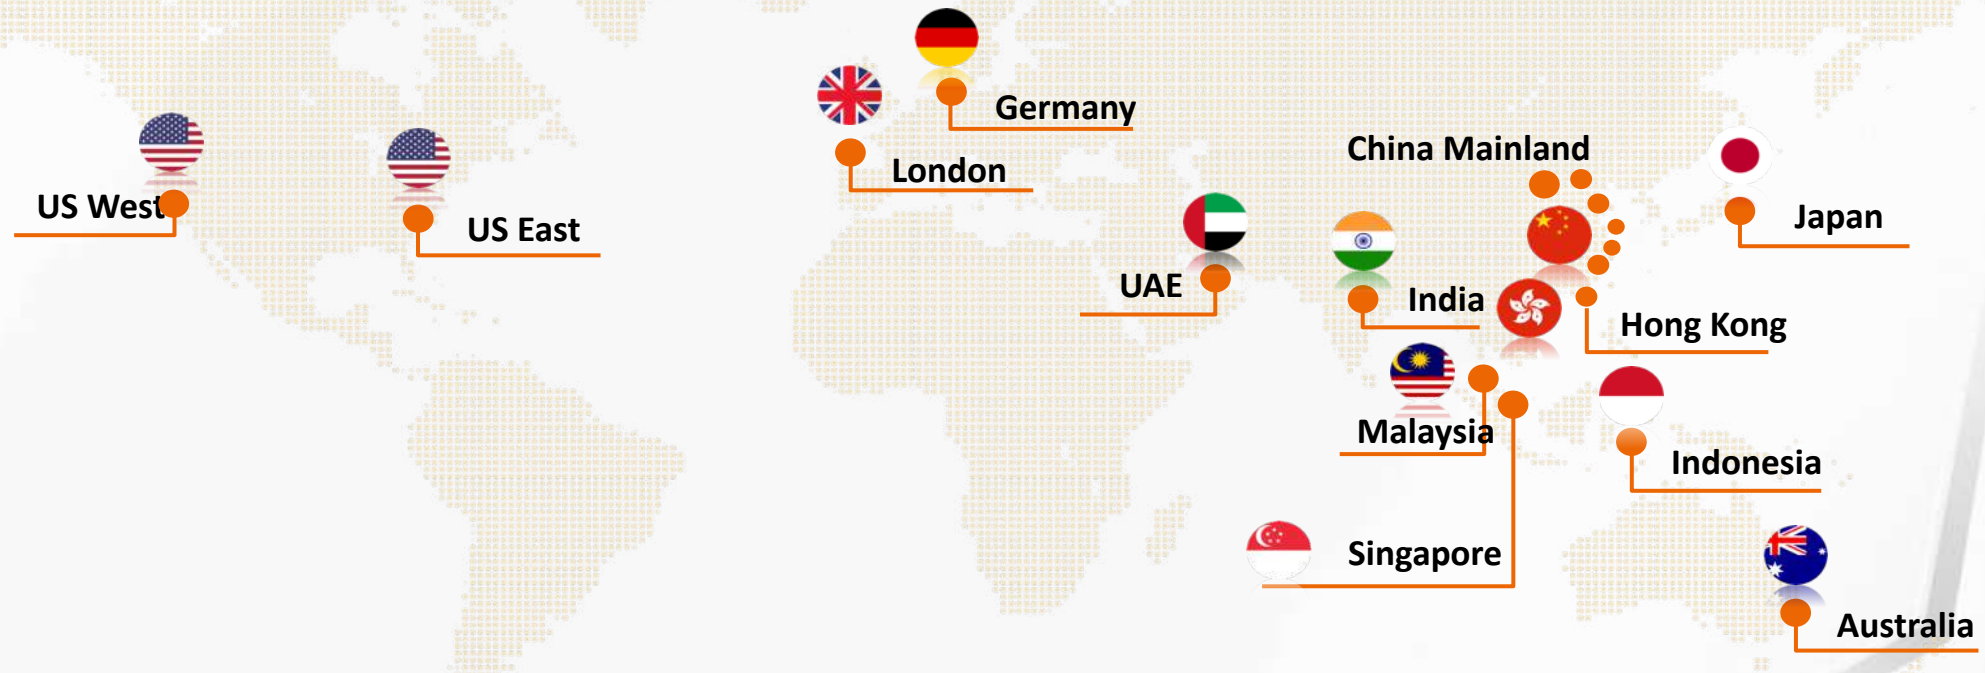
ACK—Alibaba cloud Container service for Kubernetes





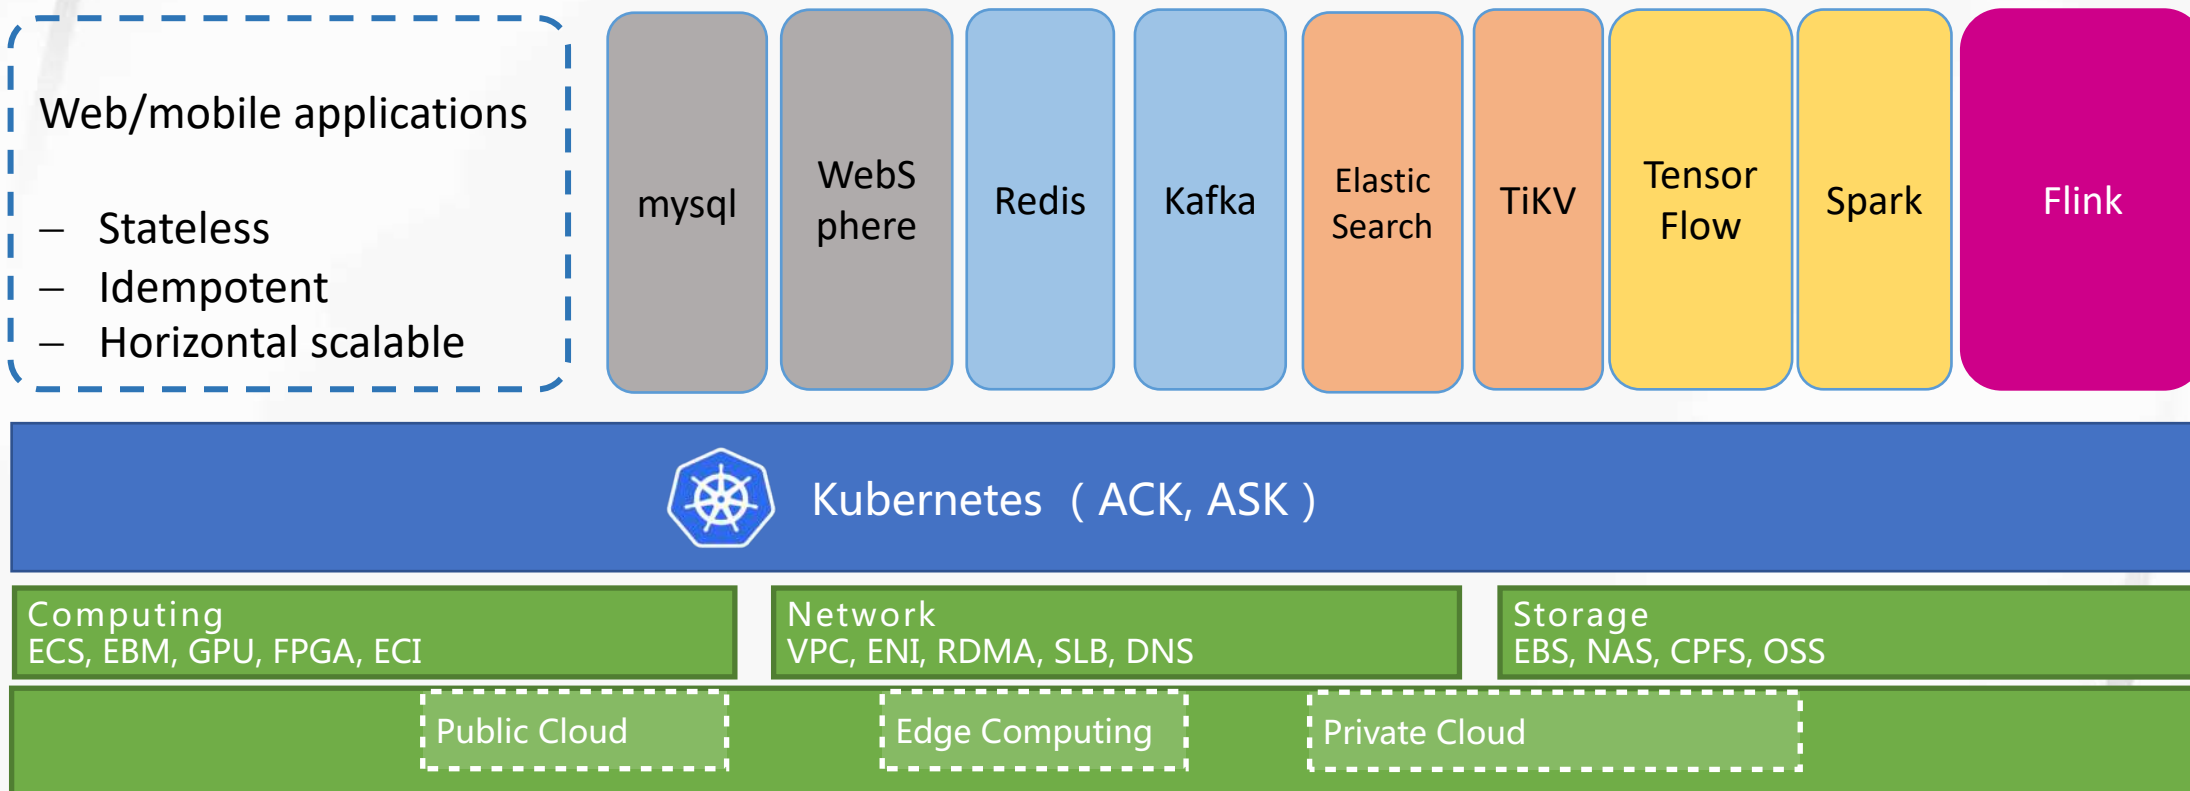
ACK is available in 18 Regions

- IDC: market share of Alibaba Cloud is 43%, ranks top 1 in China
- Gartner: Alibaba Cloud led the Asia Pacific market for IaaS and IUS with 19.6% market share





More and more workloads run on Kubernetes





02

Flink on K8s



Flink on K8s – Standalone

jobmanager-deployment.yaml

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: flink-jobmanager
spec:
  replicas: 1
  template:
    metadata:
      labels:
        app: flink
        component: jobmanager
    spec:
      containers:
        - name: jobmanager
          image: flink:latest
          args:
            - jobmanager
```

taskmanager-deployment.yaml

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: flink-taskmanager
spec:
  replicas: 2
  template:
    metadata:
      labels:
        app: flink
        component: taskmanager
    spec:
      containers:
        - name: taskmanager
          image: flink:latest
          args:
            - taskmanager
```

jobmanager-service.yaml

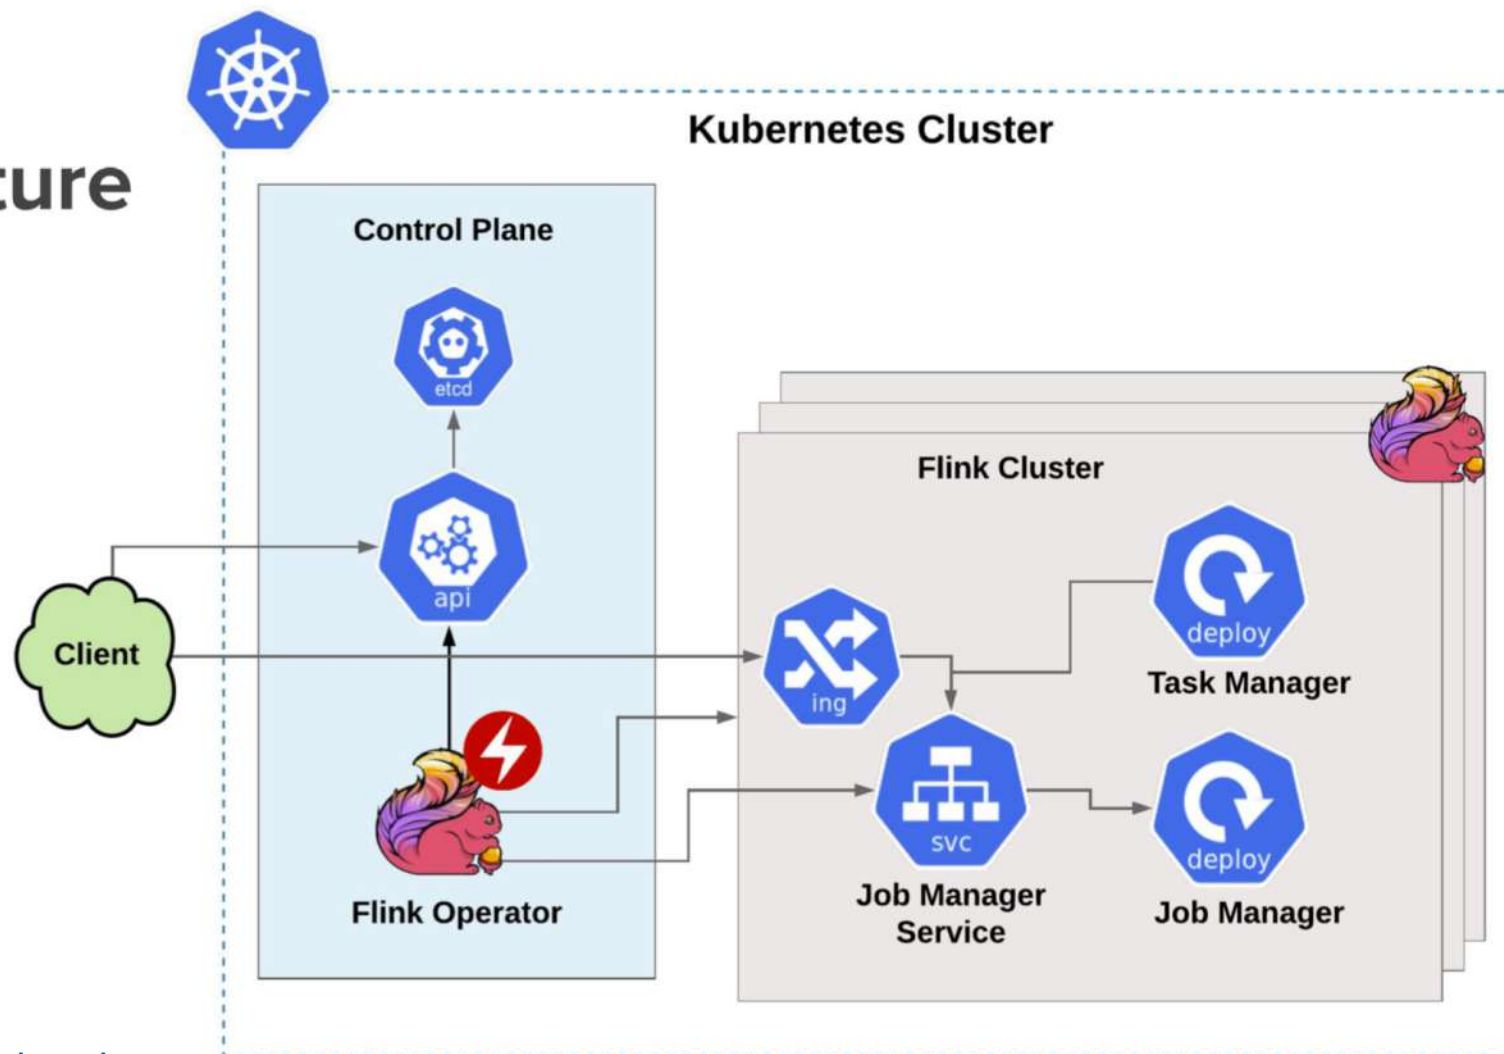
```
apiVersion: v1
kind: Service
metadata:
  name: flink-jobmanager
spec:
  ports:
    - name: rpc
      port: 6123
    - name: blob
      port: 6124
    - name: query
      port: 6125
    - name: ui
      port: 8081
  selector:
    app: flink
    component: jobmanager
```

Helm Chart: <https://github.com/docker-flink/examples>



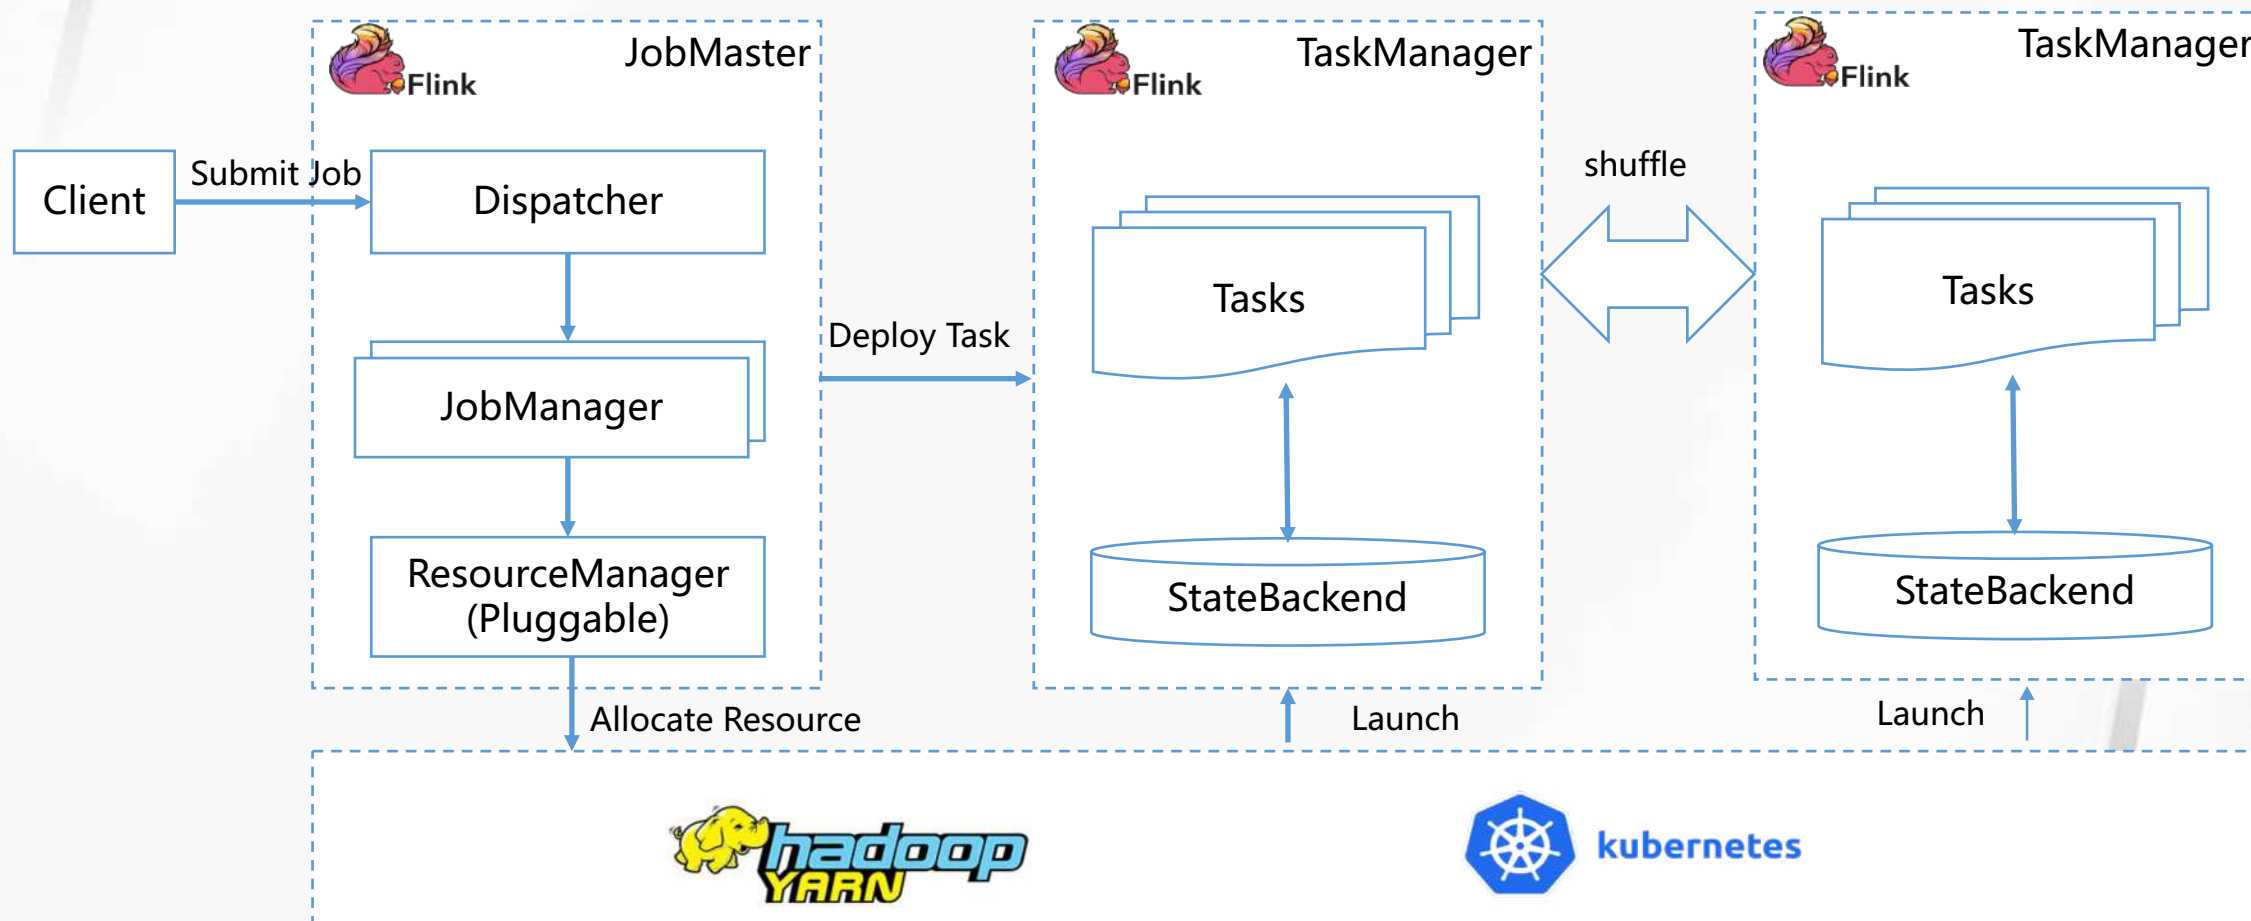
Flink on K8s – Operator

Solution Architecture



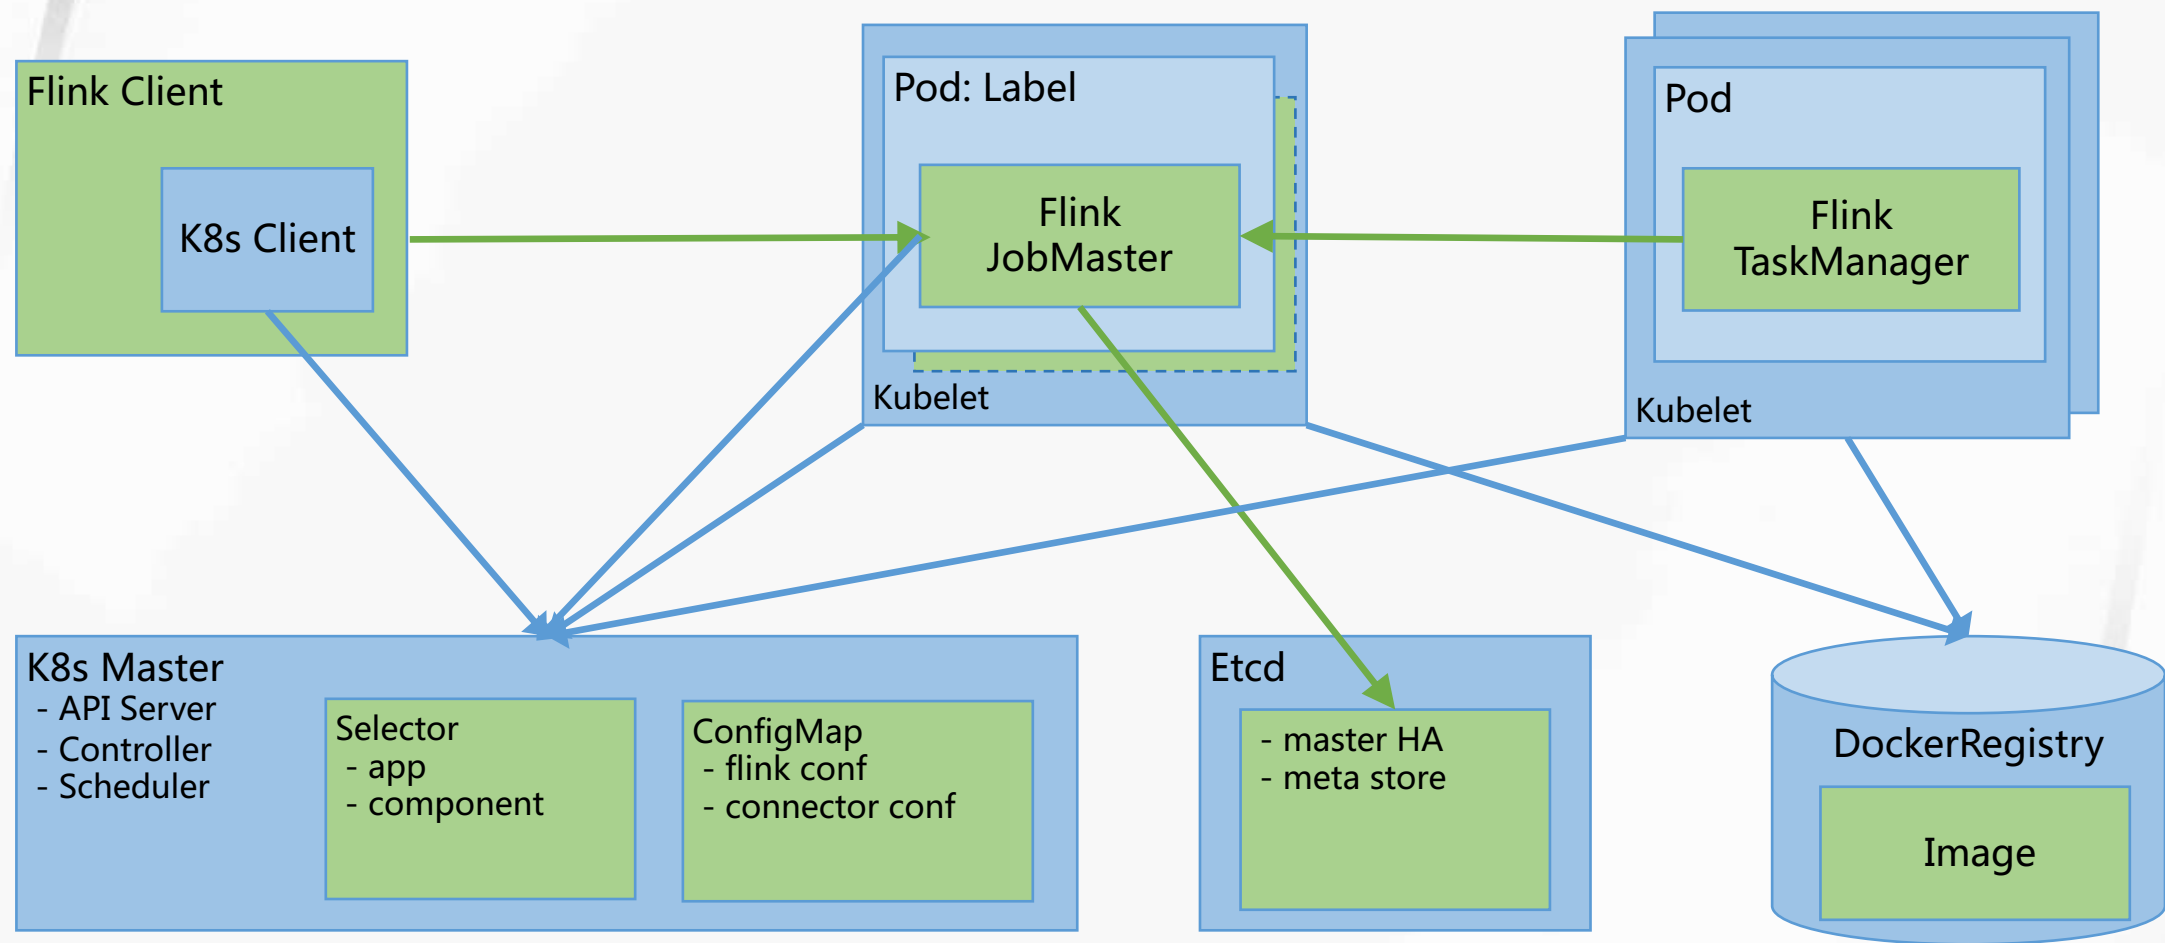


Flink on K8s – Natively





Flink on K8s – Natively



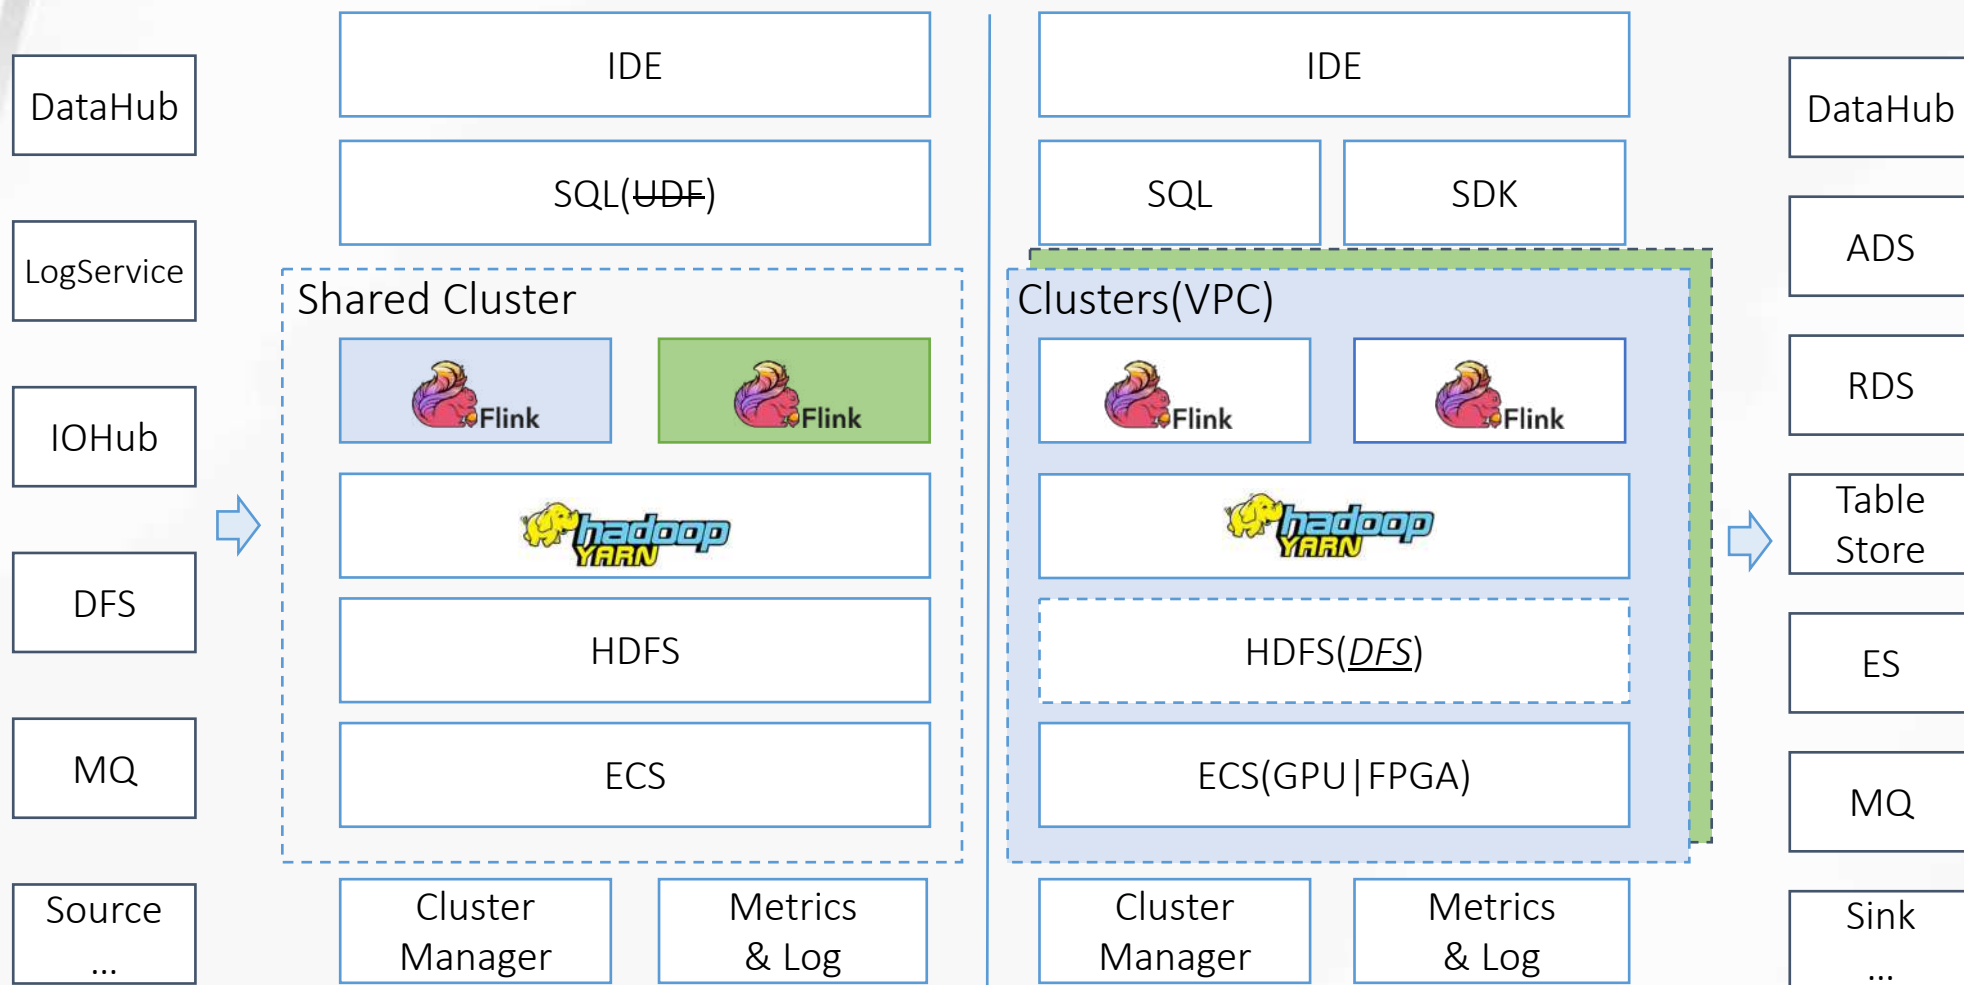


03

大数据平台

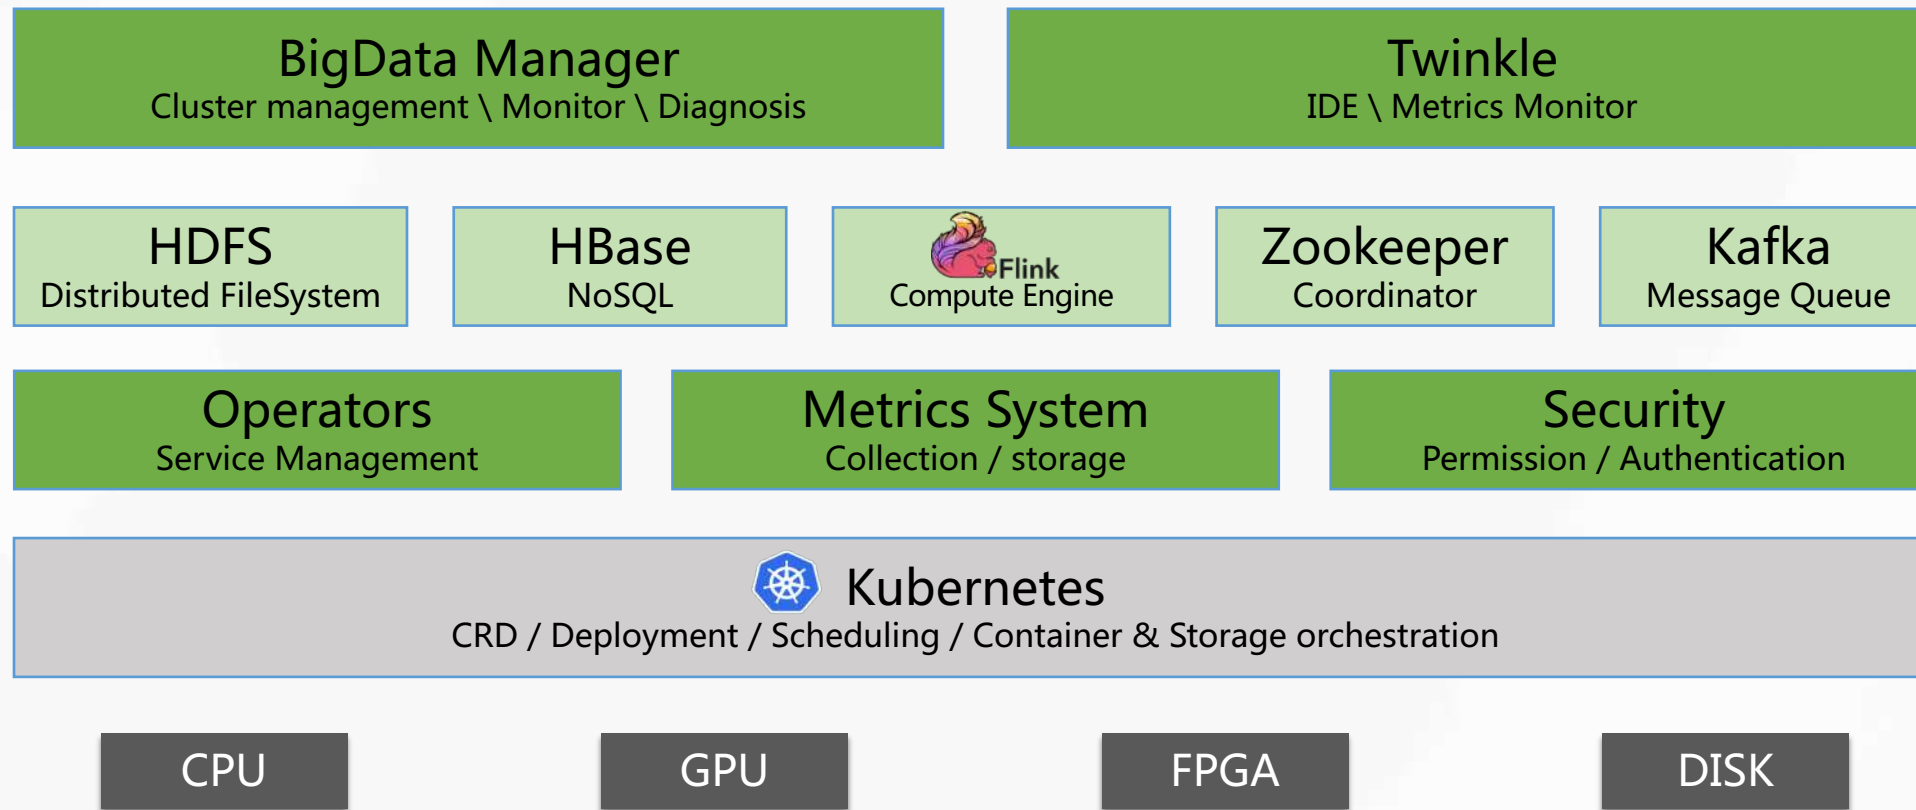


大数据平台 — on YARN



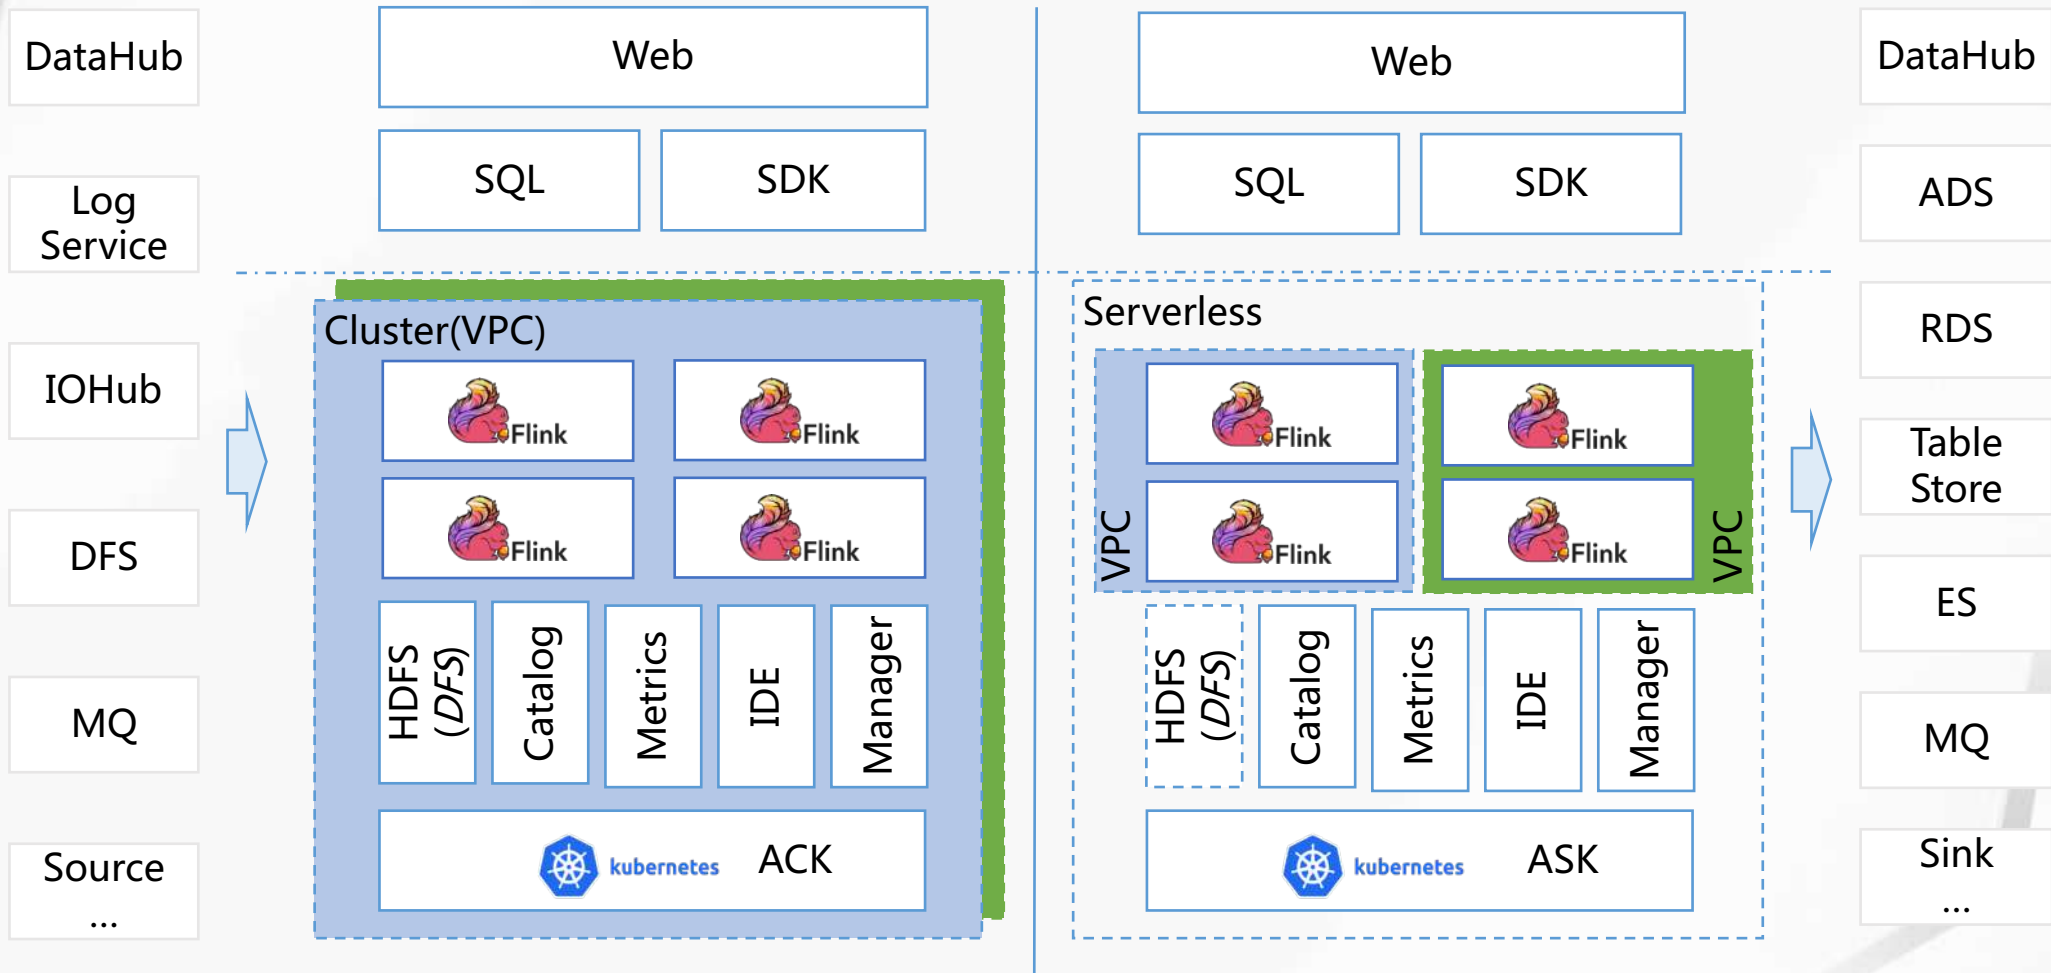


大数据平台 — on Kubernetes





大数据平台



大数据平台



Apache Flink

100%

大数据平台



Apache Flink

Apsara Bigdata Manager

运维 监控 管理 admin

服务 节点

服务 节点

服务

hdfs 操作 blink

概览 监控大盘 状态 实例

Zookeeper

Kafka

Hdfs

Hbase

Twinkle

MetricsGateway

MetricsImporter

Influxdb

服务汇总

属性 内容

create_time	2019-06-20 19:49:52
hadoop_image_tag	2.8.0-adp2.3.9
hfs_nameservice	kubecon
data_node_limit_cpu	1
hdfs_cluster_name	hdfs-cluster
data_node_limit_memory	2048Mi
journal_node_request_memory	2048Mi
pv_storage_name	fast-disks
journal_node_limit_memory	2048Mi
pv_size	90Gi

共 21 条

CPU allocation (cores)

req_cpu 10

limit_cpu 10

Memory allocation (Gb)

req_mem 18

limit_mem 18

2019-06-21 04:20 ~ 2019-06-21 10:20

刷新

打开自动刷新

OVERVIEW

Live DataNodes

6

file数量

Dead DataNodes

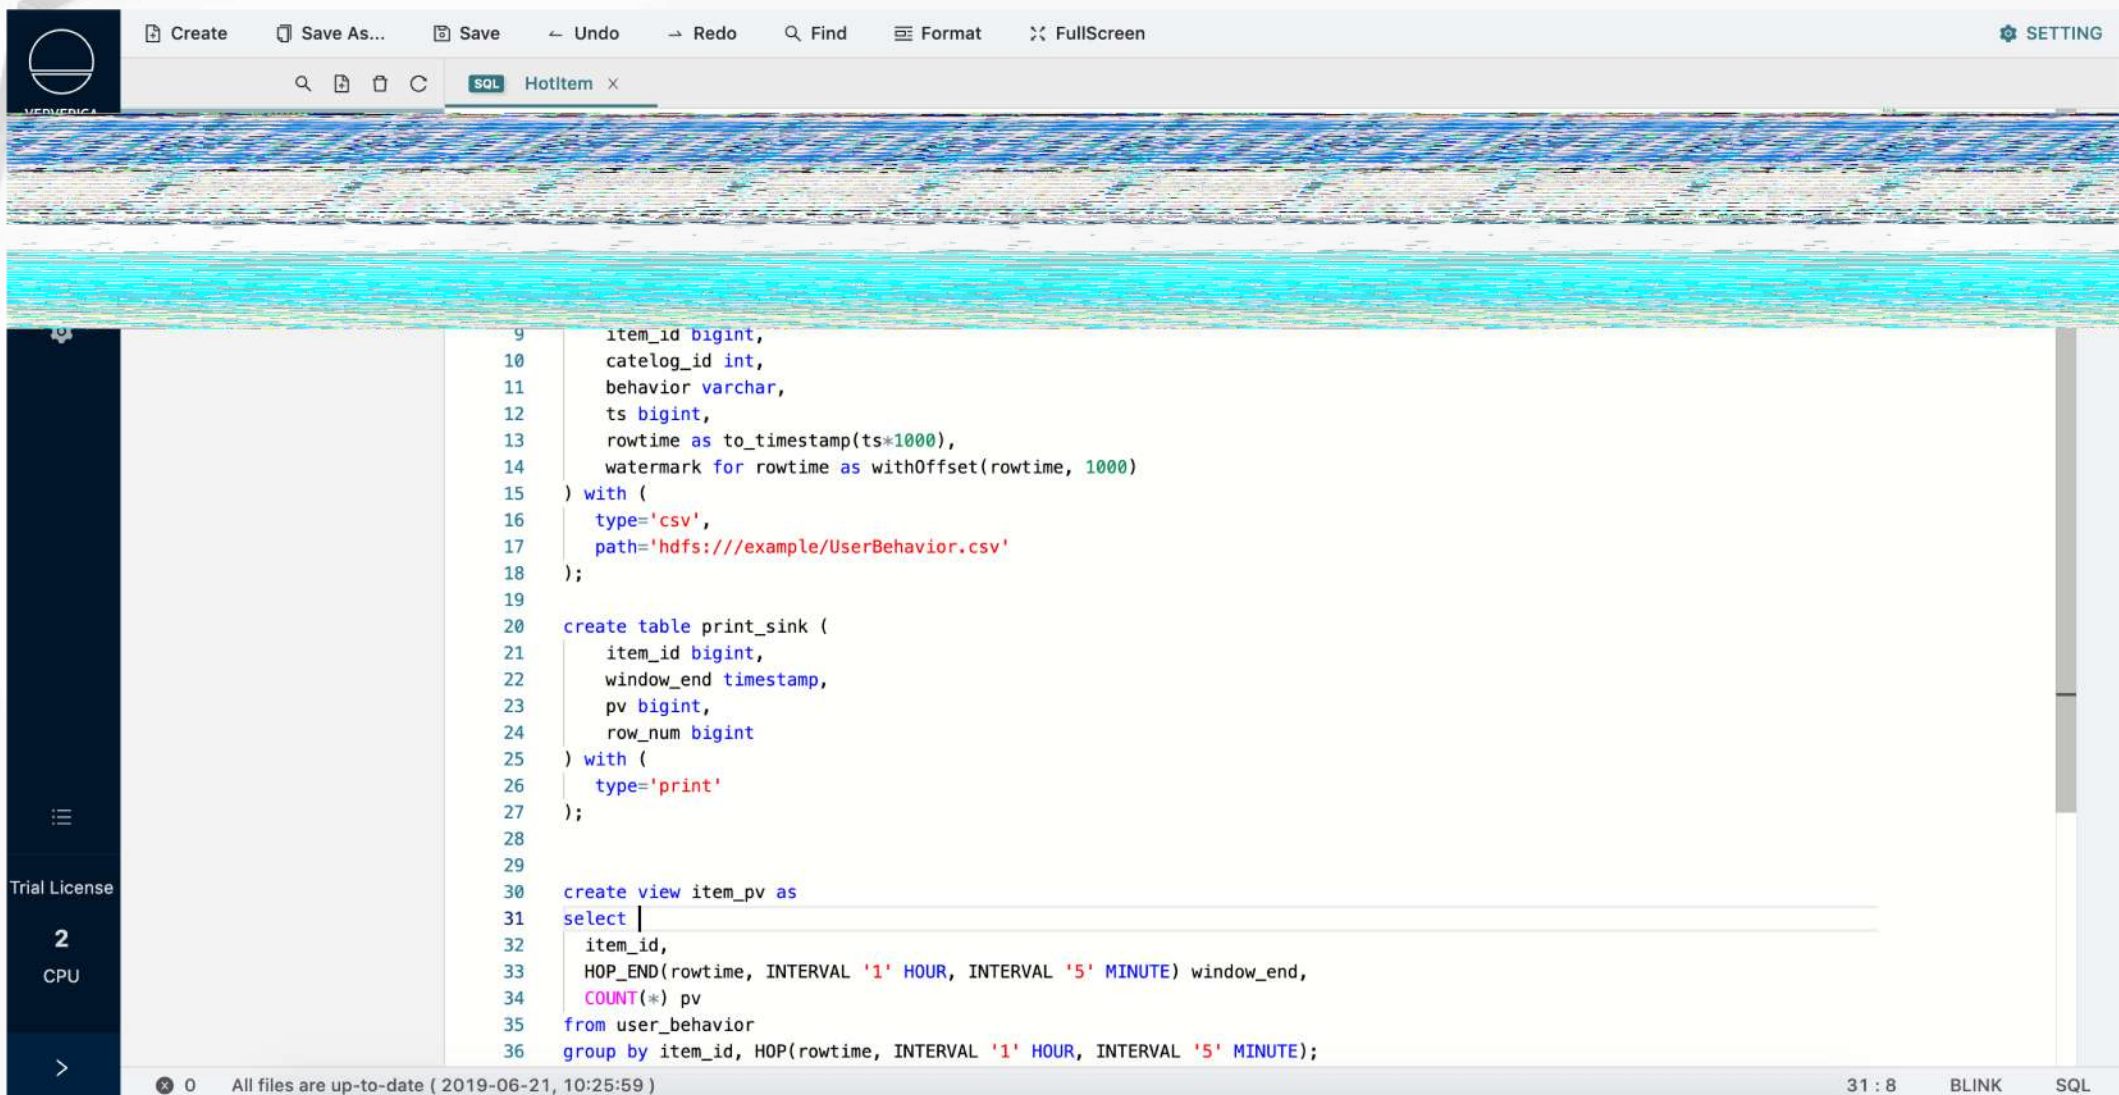
0

Block数量

大数据平台



Apache Flink





04

展望



展望

- Batch on K8s
 - Scheduling Throughput
 - Isolation
 - Shuffle Service
- Hybrid Workloads on K8s
 - Online Service
 - Streaming Job
 - Batch Job



Apache Flink

谢谢！
Q&A

【2 群】 Apache Flink C...



扫一扫群二维码，立刻加入该群。

How Kubernetes works?

